

Throughput-Quality Trade-off on OpenWebText
anchored at lossy $l=1.0$ (= strict speculative decoding)

