

Throughput-Quality Trade-off on OpenWebText
 anchored at lossy $l=1.0$ (= Leviathan strict speculative decoding)

